

# TensorFlow: Federated Learning for Image Classification

Andrea Hofer, Ignatz Ötzlinger, Sabine Hasenleithner,  
Kim Schmider



---

image source:

[https://commons.wikimedia.org/wiki/File:Tensorflow\\_logo.svg](https://commons.wikimedia.org/wiki/File:Tensorflow_logo.svg)

# Introduction

- ▶ Introduction to Machine Learning and Deep Learning
- ▶ Image Recognition
- ▶ CIFAR-10 dataset
- ▶ Practical example
- ▶ Why Federated Learning?
- ▶ FederatedAveraging Algorithm

# What is Machine Learning?

- ▶ Big Data: There is no data like more data

# What is Machine Learning?

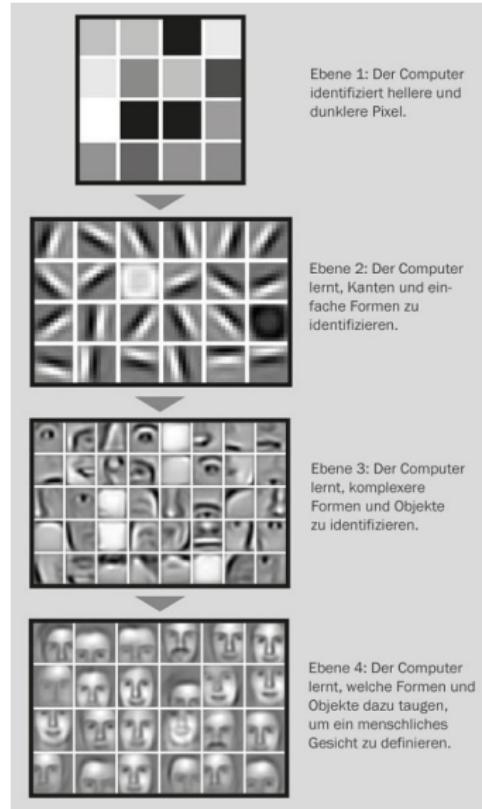
- ▶ Big Data: There is no data like more data
- ▶ Deep Learning: Error rate below 4%

# What is Machine Learning?

- ▶ Big Data: There is no data like more data
- ▶ Deep Learning: Error rate below 4%



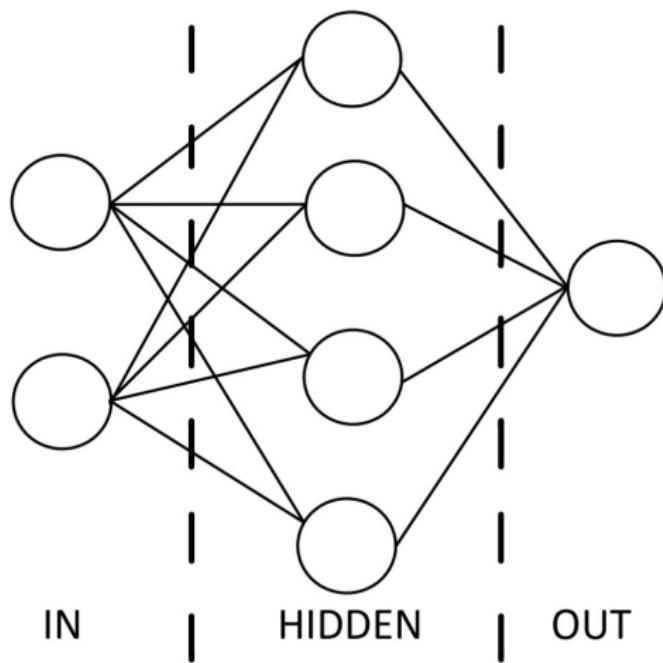
source: <https://www.computerwoche.de/a/was-sie-ueber-maschinelles-lernen-wissen-muessen,3329560>



source:

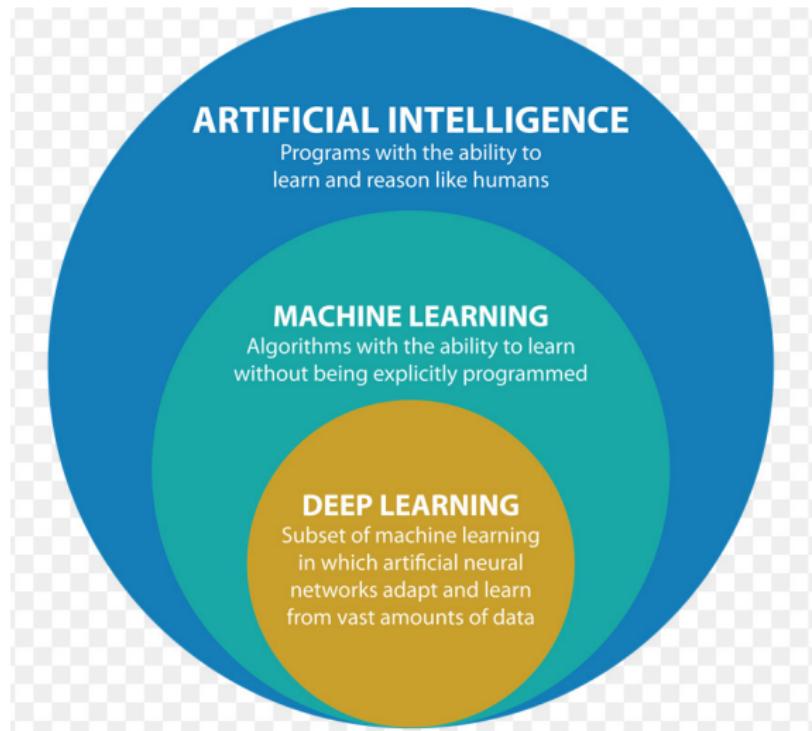
[http://www.spektrum.de/news/maschinenlernen-deep-learning  
-macht-kuenstliche-intelligenz-praxistauglich/1220451](http://www.spektrum.de/news/maschinenlernen-deep-learning-macht-kuenstliche-intelligenz-praxistauglich/1220451)

# Neural Network



source: own figure

# Deep Learning



source: <https://de.cleanpng.com/png-vy5m68>

# Image Recognition

## ► Classification

# Image Recognition

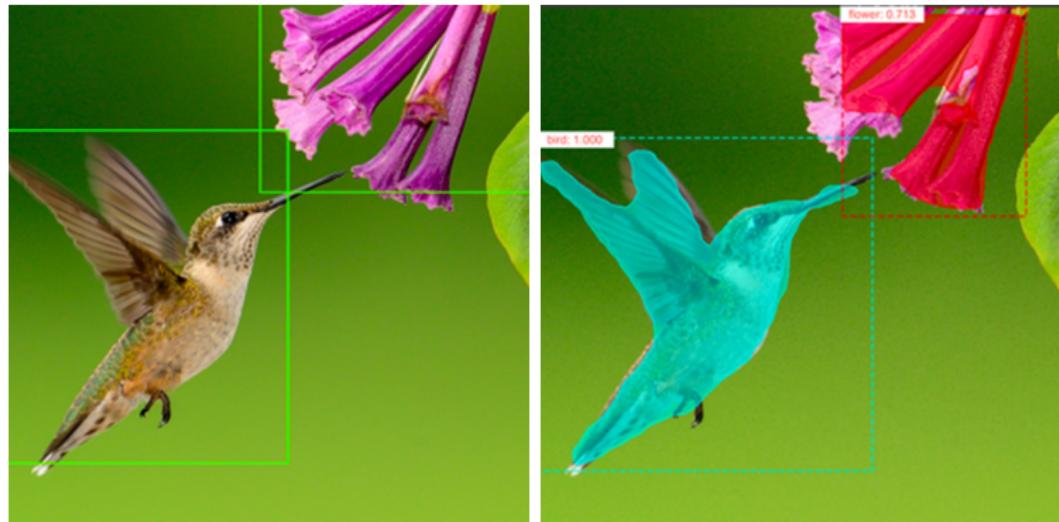
- ▶ Classification
- ▶ Object recognition

# Image Recognition

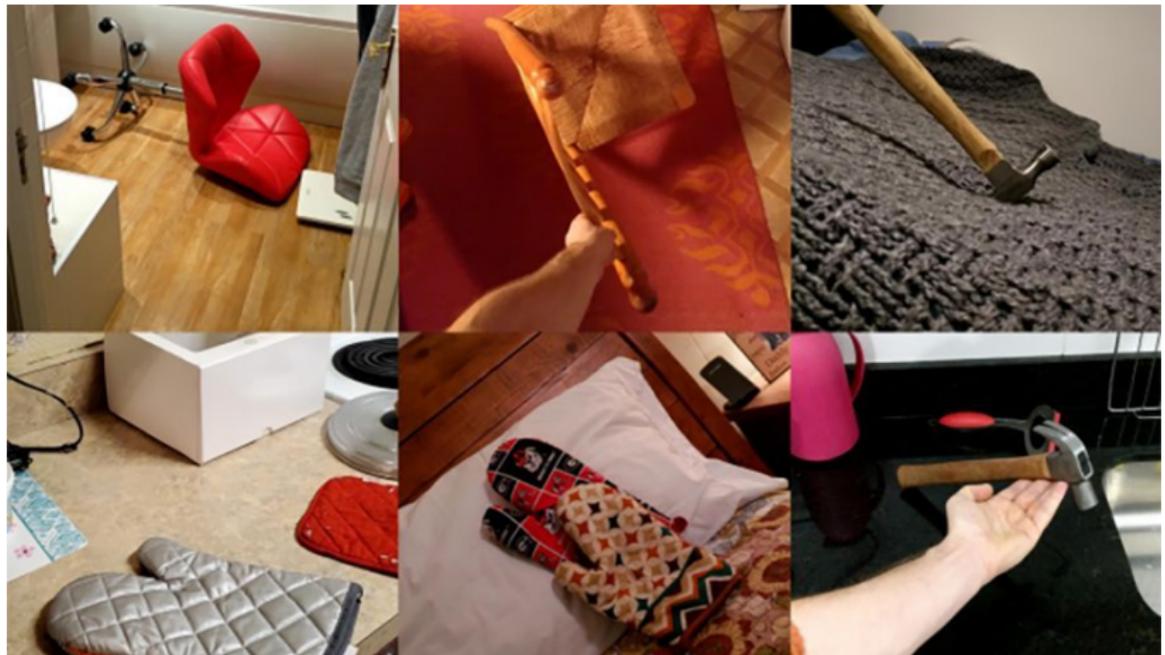
- ▶ Classification
- ▶ Object recognition
- ▶ Instance segmentation

# Image Recognition

- ▶ Classification
- ▶ Object recognition
- ▶ Instance segmentation



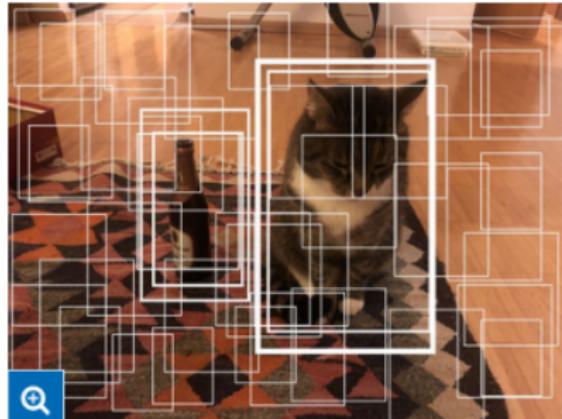
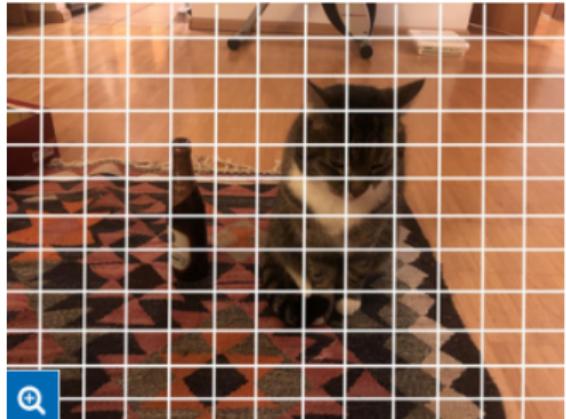
source: <https://www.burda.com/de/news/tech-einmaleins-bild-erkennung-und-ki>



source:

<https://www.spiegel.de/netzwelt/web/bilderkennung-an-diese-n-fotos-scheitert-kuenstliche-intelligenz-a-1301234.html>

# YOLO Algorithm



source: <https://www.sigs-datacom.de/trendletter/2018-10/4-wie-man-in-echtzeit-mehrere-objekte-mit-deep-learning-und-yolo-erkennen-und-klassifizieren-kann.html>

# CIFAR-10

- ▶ CIFAR-10 dataset (Canadian Institute For Advanced Research)

## CIFAR-10

- ▶ CIFAR-10 dataset (Canadian Institute For Advanced Research)
- ▶ Collection of images

## CIFAR-10

- ▶ CIFAR-10 dataset (Canadian Institute For Advanced Research)
- ▶ Collection of images
- ▶ Contains 60.000 images

# CIFAR-10

- ▶ CIFAR-10 dataset (Canadian Institute For Advanced Research)
- ▶ Collection of images
- ▶ Contains 60.000 images
- ▶ 10 different classes:
  1. airplanes,
  2. cars,
  3. birds,
  4. cats,
  5. deer,
  6. dogs,
  7. frogs,
  8. horses,
  9. ships,
  10. trucks.

# CIFAR-10

- ▶ CIFAR-10 dataset (Canadian Institute For Advanced Research)
- ▶ Collection of images
- ▶ Contains 60.000 images
- ▶ 10 different classes:
  1. airplanes,
  2. cars,
  3. birds,
  4. cats,
  5. deer,
  6. dogs,
  7. frogs,
  8. horses,
  9. ships,
  10. trucks.
- ▶ 6.000 images per class

# CIFAR-10

- ▶ CIFAR-10 dataset (Canadian Institute For Advanced Research)
- ▶ Collection of images
- ▶ Contains 60.000 images
- ▶ 10 different classes:
  1. airplanes,
  2. cars,
  3. birds,
  4. cats,
  5. deer,
  6. dogs,
  7. frogs,
  8. horses,
  9. ships,
  10. trucks.
- ▶ 6.000 images per class
- ▶ Size 32 x 32

# Image Classification

How does the image classification work?

<https://www.tensorflow.org/tutorials/images/cnn>

This tutorial demonstrates training a simple Convolutional Neural Network (CNN) to classify CIFAR images. (TensorFlow)

## SGD – Stochastic gradient descent

- ▶ Iterative method for optimizing an objective function

## SGD – Stochastic gradient descent

- ▶ Iterative method for optimizing an objective function
- ▶ Replaces the actual gradient (calculated from the entire data set) by an estimate thereof (calculated from a randomly selected subset of the data)

## SGD – Stochastic gradient descent

- ▶ Iterative method for optimizing an objective function
- ▶ Replaces the actual gradient (calculated from the entire data set) by an estimate thereof (calculated from a randomly selected subset of the data)
- ▶ True gradient of  $Q(w)$  is approximated by a gradient at a single example

## SGD – Stochastic gradient descent

- ▶ Iterative method for optimizing an objective function
- ▶ Replaces the actual gradient (calculated from the entire data set) by an estimate thereof (calculated from a randomly selected subset of the data)
- ▶ True gradient of  $Q(w)$  is approximated by a gradient at a single example
- ▶  $w := w - \eta \nabla Q_i(w)$

## SGD – Stochastic gradient descent

- ▶ Iterative method for optimizing an objective function
- ▶ Replaces the actual gradient (calculated from the entire data set) by an estimate thereof (calculated from a randomly selected subset of the data)
- ▶ True gradient of  $Q(w)$  is approximated by a gradient at a single example
- ▶  $w := w - \eta \nabla Q_i(w)$
- ▶ Algorithm sweeps through the training set and performs the above update for each training example. Several passes can be made over the training set until the algorithm converges. If this is done, the data can be shuffled for each pass to prevent cycles

## SGD – Possible pseudocode

choose an initial vector of parameters  $w$  and learning rate  $\eta$

## SGD – Possible pseudocode

choose an initial vector of parameters  $w$  and learning rate  $\eta$   
**repeat** until an approximate minimum is obtained **do**

## SGD – Possible pseudocode

choose an initial vector of parameters  $w$  and learning rate  $\eta$   
**repeat** until an approximate minimum is obtained **do**  
    randomly shuffle examples in the training set

## SGD – Possible pseudocode

```
choose an initial vector of parameters  $w$  and learning  
rate  $\eta$   
repeat until an approximate minimum is obtained do  
    randomly shuffle examples in the training set  
    for  $i = 1, 2, \dots, n$  do
```

## SGD – Possible pseudocode

choose an initial vector of parameters  $w$  and learning rate  $\eta$

**repeat** until an approximate minimum is obtained **do**

randomly shuffle examples in the training set

**for**  $i = 1, 2, \dots, n$  **do**

$w := w - \eta \nabla Q_i(w)$

## SGD – Possible pseudocode

choose an initial vector of parameters  $w$  and learning rate  $\eta$

**repeat** until an approximate minimum is obtained **do**

randomly shuffle examples in the training set

**for**  $i = 1, 2, \dots, n$  **do**

$w := w - \eta \nabla Q_i(w)$

source: Bottou, Léon (1998). "Online Algorithms and Stochastic Approximations". Online Learning and Neural Networks. Cambridge University Press. ISBN 978-0-521-65263-6.

# Why Federated Learning?

# Federated Learning

- ▶ Calculations on end device

## Federated Learning

- ▶ Calculations on end device
- ▶ Results of calculations will be sent to global model

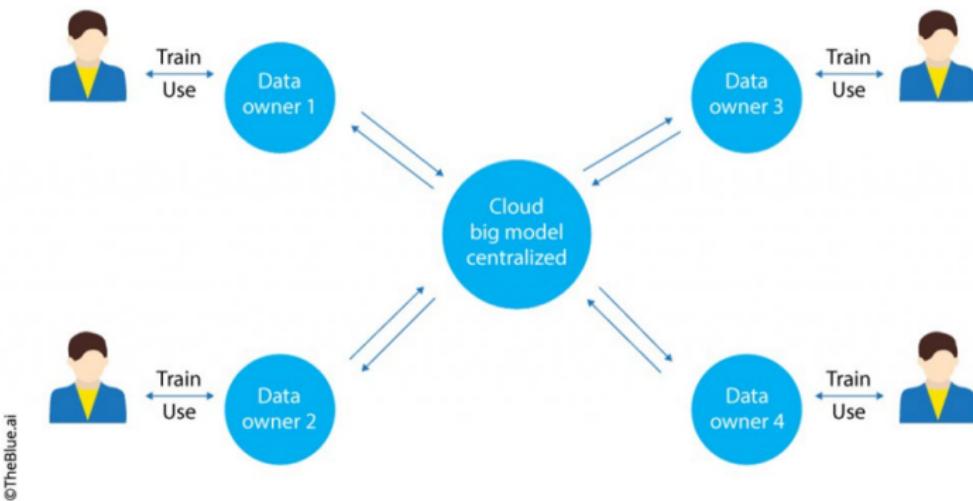
## Federated Learning

- ▶ Calculations on end device
- ▶ Results of calculations will be sent to global model
- ▶ Data privacy

# Federated Learning

- ▶ Calculations on end device
- ▶ Results of calculations will be sent to global model
- ▶ Data privacy

Federated Learning



©TheBlue.ai

source: <https://theblue.ai/blog-de/federated-learning-federales-lernen>

## Federated Learning for image classification

- ▶ Can be simulated at TensorFlow:

[https://www.tensorflow.org/federated/tutorials/federated\\_learning\\_for\\_image\\_classification](https://www.tensorflow.org/federated/tutorials/federated_learning_for_image_classification)

## Federated Learning for image classification

- ▶ Can be simulated at TensorFlow:  
[https://www.tensorflow.org/federated/tutorials/federated\\_learning\\_for\\_image\\_classification](https://www.tensorflow.org/federated/tutorials/federated_learning_for_image_classification)
- ▶ this simulation example uses the NIST (National Institute of Standards and Technology) dataset

# NIST Handprinted Forms and Characters Database

- ▶ 3600 writers

# NIST Handprinted Forms and Characters Database

- ▶ 3600 writers
- ▶ 810,000 character images

# NIST Handprinted Forms and Characters Database

- ▶ 3600 writers
- ▶ 810,000 character images

HANDWRITING SAMPLE FORM

NAME	DATE	CITY	STATE ZIP
[REDACTED]	8-3-89	MUNEN CITY	MI 48452
This sample of handwriting is being collected for use in testing computer recognition of hand printed numbers and letters. Please print the following characters in the boxes that appear below:			
0 1 2 3 4 5 6 7 8 9	0 1 2 3 4 5 6 7 8 9	0 1 2 3 4 5 6 7 8 9	0 1 2 3 4 5 6 7 8 9
67	701	3752	80759
158	4586	32123	832656
7481	80530	419219	67
61738	729658	75	390
109334	40	675	4234
			46002

gyxlaapkdebtsirumwfqjenhocv  
9yXlaKpd5bTzjruuWF9jenhocv  
ZXSBNGECMYWQTKFLUOHPIRVDJA  
ZXSBNGECMYWQTKFLUOHPIRVDJA

Please print the following text in the box below:  
We, the People of the United States, in order to form a more perfect Union, establish Justice, insure domestic Tranquility, provide for the common Defense, promote the general Welfare, and secure the Blessings of Liberty to ourselves and our posterity, do ordain and establish this CONSTITUTION for the United States of America.

We, the people of the United States, in order to form a more perfect Union, establish Justice, insure domestic Tranquility, provide for the common Defense, promote the general Welfare, and secure the Blessings of liberty to ourselves and our posterity, do ordain and establish this CONSTITUTION for the United States of America.

source: <https://www.nist.gov/srd/nist-special-database-19>

# Short overview Simulation results

```
>>> state, metrics = iterative_process.next(state, federated_train_data)
>>> print('round 1, metrics={}'.format(metrics))
round 1, metrics=<sparse_categorical_accuracy=0.14228394627571106,loss=2.9735283851623535>
>>> for round_num in range(2, 11):
...     state, metrics = iterative_process.next(state, federated_train_data)
...     print('round {:2d}, metrics={}'.format(round_num, metrics))
...
round 2, metrics=<sparse_categorical_accuracy=0.18034979701042175,loss=2.6374611854553223>
round 3, metrics=<sparse_categorical_accuracy=0.21903292834758759,loss=2.547621965408325>
round 4, metrics=<sparse_categorical_accuracy=0.2531892955303192,loss=2.378369092941284>
round 5, metrics=<sparse_categorical_accuracy=0.32098764181137085,loss=2.047738552093506>
round 6, metrics=<sparse_categorical_accuracy=0.36450618505477905,loss=1.9402925968170166>
round 7, metrics=<sparse_categorical_accuracy=0.419444417953491,loss=1.7680468559265137>
round 8, metrics=<sparse_categorical_accuracy=0.4552469253540039,loss=1.6488759517669678>
round 9, metrics=<sparse_categorical_accuracy=0.529629647731781,loss=1.4977922439575195>
round 10, metrics=<sparse_categorical_accuracy=0.5630658268928528,loss=1.3877880573272705>
```

source: own figure

## Short overview Simulation results

```
>>> state, metrics = iterative_process.next(state, federated_train_data)
>>> print('round 1, metrics={}'.format(metrics))
round 1, metrics=<sparse_categorical_accuracy=0.14228394627571106,loss=2.9735283851623535>
>>> for round_num in range(2, 11):
...     state, metrics = iterative_process.next(state, federated_train_data)
...     print('round {:2d}, metrics={}'.format(round_num, metrics))
...
round 2, metrics=<sparse_categorical_accuracy=0.18034979701042175,loss=2.6374611854553223>
round 3, metrics=<sparse_categorical_accuracy=0.21903292834758759,loss=2.547621965408325>
round 4, metrics=<sparse_categorical_accuracy=0.2531892955303192,loss=2.378369092941284>
round 5, metrics=<sparse_categorical_accuracy=0.32098764181137085,loss=2.047738552093506>
round 6, metrics=<sparse_categorical_accuracy=0.36450618505477905,loss=1.9402925968170166>
round 7, metrics=<sparse_categorical_accuracy=0.4194444417953491,loss=1.7680468559265137>
round 8, metrics=<sparse_categorical_accuracy=0.4552469253540039,loss=1.6488759517669678>
round 9, metrics=<sparse_categorical_accuracy=0.529629647731781,loss=1.4977922439575195>
round 10, metrics=<sparse_categorical_accuracy=0.5630658268928528,loss=1.3877880573272705>
```

source: own figure

Training loss is decreasing after each round of federated training, indicating the model is converging

## Short overview Simulation results

```
>>> state, metrics = iterative_process.next(state, federated_train_data)
>>> print('round 1, metrics={}'.format(metrics))
round 1, metrics=<sparse_categorical_accuracy=0.14228394627571106,loss=2.9735283851623535>
>>> for round_num in range(2, 11):
...     state, metrics = iterative_process.next(state, federated_train_data)
...     print('round {:2d}, metrics={}'.format(round_num, metrics))
...
round 2, metrics=<sparse_categorical_accuracy=0.18034979701042175,loss=2.6374611854553223>
round 3, metrics=<sparse_categorical_accuracy=0.21903292834758759,loss=2.547621965408325>
round 4, metrics=<sparse_categorical_accuracy=0.2531892955303192,loss=2.378369092941284>
round 5, metrics=<sparse_categorical_accuracy=0.32098764181137085,loss=2.047738552093506>
round 6, metrics=<sparse_categorical_accuracy=0.36450618505477905,loss=1.9402925968170166>
round 7, metrics=<sparse_categorical_accuracy=0.4194444417953491,loss=1.7680468559265137>
round 8, metrics=<sparse_categorical_accuracy=0.4552469253540039,loss=1.6488759517669678>
round 9, metrics=<sparse_categorical_accuracy=0.529629647731781,loss=1.4977922439575195>
round 10, metrics=<sparse_categorical_accuracy=0.5630658268928528,loss=1.3877880573272705>
```

source: own figure

Training loss is decreasing after each round of federated training, indicating the model is converging  
The average metrics over all batches of data trained across all clients in the round

# Example Uses of Federated Learning

- ▶ Medicine

## Example Uses of Federated Learning

- ▶ Medicine
- ▶ Touch keyboard input prediction

## Example Uses of Federated Learning

- ▶ Medicine
- ▶ Touch keyboard input prediction
- ▶ With learning based on user interactions labels are directly available

# Practical Issues With Federated Learning

- ▶ Datasets are not representative of population

# Practical Issues With Federated Learning

- ▶ Datasets are not representative of population
- ▶ Data differs from traditional centralized sources

# Practical Issues With Federated Learning

- ▶ Datasets are not representative of population
- ▶ Data differs from traditional centralized sources
- ▶ Users' usage styles differ greatly

# Practical Issues With Federated Learning

- ▶ Datasets are not representative of population
- ▶ Data differs from traditional centralized sources
- ▶ Users' usage styles differ greatly
- ▶ Many clients each with small datasets

# Practical Issues With Federated Learning

- ▶ Datasets are not representative of population
- ▶ Data differs from traditional centralized sources
- ▶ Users' usage styles differ greatly
- ▶ Many clients each with small datasets
- ▶ Mobile devices are often offline or on a slow or expensive connection

# Practical Issues With Federated Learning

- ▶ Datasets are not representative of population
- ▶ Data differs from traditional centralized sources
- ▶ Users' usage styles differ greatly
- ▶ Many clients each with small datasets
- ▶ Mobile devices are often offline or on a slow or expensive connection
- ▶ Time zone differences

# Resource Utilization

- ▶ Data center: Computational costs dominate

## Resource Utilization

- ▶ Data center: Computational costs dominate
- ▶ Federated Learning: Communication costs dominate, computation essentially free

## Resource Utilization

- ▶ Data center: Computational costs dominate
- ▶ Federated Learning: Communication costs dominate, computation essentially free
- ▶ Number of updates is minimized:

## Resource Utilization

- ▶ Data center: Computational costs dominate
- ▶ Federated Learning: Communication costs dominate, computation essentially free
- ▶ Number of updates is minimized:
  - ▶ by running on more clients in parallel (diminishing returns), or

## Resource Utilization

- ▶ Data center: Computational costs dominate
- ▶ Federated Learning: Communication costs dominate, computation essentially free
- ▶ Number of updates is minimized:
  - ▶ by running on more clients in parallel (diminishing returns), or
  - ▶ by performing more computations between updates.

# FederatedAveraging Algorithm

*Server*

```
Main() do
    initialize  $w_0$ 
    for each round  $t = 1, 2, \dots$  do
         $S_t \leftarrow$  random fraction of clients
        for each client  $k \in S_t$  in parallel do
             $(w_{t+1}^k, n_k) \leftarrow \text{ClientUpdate}(k, w_t)$ 
```

# FederatedAveraging Algorithm

*Server*

```
Main() do
    initialize  $w_0$ 
    for each round  $t = 1, 2, \dots$  do
         $S_t \leftarrow$  random fraction of clients
        for each client  $k \in S_t$  in parallel do
             $(w_{t+1}^k, n_k) \leftarrow \text{ClientUpdate}(k, w_t)$ 
```

*Client k*

```
ClientUpdate( $k, w$ ) do
    for each epoch do
         $w \leftarrow w - \eta \nabla Q_i(w)$ 
    return  $(w, n_k)$  to server
```

# FederatedAveraging Algorithm

*Server*

```
Main() do
    initialize  $w_0$ 
    for each round  $t = 1, 2, \dots$  do
         $S_t \leftarrow$  random fraction of clients
        for each client  $k \in S_t$  in parallel do
             $(w_{t+1}^k, n_k) \leftarrow \text{ClientUpdate}(k, w_t)$ 
         $n \leftarrow n_1 + n_2 + \dots + n_k$ 
         $w_{t+1} \leftarrow \sum_{k=1}^K \frac{n_k}{n} w_{t+1}^k$ 
```

*Client k*

```
ClientUpdate( $k, w$ ) do
    for each epoch do
         $w \leftarrow w - \eta \nabla Q_i(w)$ 
    return  $(w, n_k)$  to server
```

# FederatedAveraging Algorithm

*Server*

```
Main() do
    initialize  $w_0$ 
    for each round  $t = 1, 2, \dots$  do
         $S_t \leftarrow$  random fraction of clients
        for each client  $k \in S_t$  in parallel do
             $(w_{t+1}^k, n_k) \leftarrow \text{ClientUpdate}(k, w_t)$ 
         $n \leftarrow n_1 + n_2 + \dots + n_k$ 
         $w_{t+1} \leftarrow \sum_{k=1}^K \frac{n_k}{n} w_{t+1}^k$ 
```

*Client k*

```
ClientUpdate( $k, w$ ) do
    for each epoch do
         $w \leftarrow w - \eta \nabla Q_i(w)$ 
    return  $(w, n_k)$  to server
```

source: H. Brendan McMahan, et. al. (2016). "Communication-Efficient Learning of Deep Networks from Decentralized Data". arXiv.org.

# The End